

Neural Networks Workshop: Training and Stochastic Gradient Descent

William Guss
wguss@berkeley.edu

Phillip Kuznetsov
philkuz@berkeley.edu

University of California, Berkeley
Robotics @ Berkeley

December 1, 2015

Today we use and train Feed-Forward Artificial Neural Networks

1 Feed-Forward Neural Networks

- How It Works
- An Implementation

2 Training

- Nonconvex Optimization
- Training Demo

3 Deep Learning

Perceptron Review

Neural
Networks Pt.
2


W. Guss &
P. Kuznetsov

Feed-Forward
Neural
Networks

How It Works
An
Implementation

Training
Nonconvex
Optimization
Training Demo

Deep Learning



- Perceptrons are neural computation units which make *weighted* decisions:

$$p(\mathbf{x}) = \begin{cases} 1 & \text{if } \sum w_i x_i + b \geq 0 \\ 0 & \text{otherwise} \end{cases}$$
$$= \text{step} \left(\sum w_i x_i + b \right)$$

- Single perceptrons are not powerful enough, as seen last time with XOR.
- What if we want real valued output for tasks like predicting the temperature or stock prices?

Feedforward Neural Networks

Neural

Networks Pt. 2

W. Guss &
P. Kuznetsov


Feed-Forward Neural Networks

How It Works
An
Implementation

Training

Nonconvex
Optimization
Training Demo

Deep Learning



- Feedforward Artificial Neural Networks (ANNs) are the *continuous* extensions of perceptrons.
- ANNs can have many layers and different nodes which are *fully connected*.
- The intuition behind this model is that each neuron in the network makes a weighted decision like the perceptron. Many *stacked* decisions allows for extremely complex logic.

A Bit of Notation


Neural
Networks Pt.
2

W. Guss &
P. Kuznetsov

Feed-Forward
Neural
Networks
How It Works
An
Implementation

Training
Nonconvex
Optimization
Training Demo

Deep Learning




Definition

A **weight** on the ℓ th layer between the j th neuron on that layer and the i th neuron on the next layer is denoted $w_{ij}^\ell \in \mathbb{R}$.

Definition

The **input** to the neural network is a vector $\mathbf{x} \in \mathbb{R}^n$ and the **output** is a vector $\mathbf{y} \in \mathbb{R}^m$.

A Bit of Notation



Definition

The **output** of the i th neuron on the $(\ell + 1)$ th layer is given by the weighted sum of its inputs

$$a_i^{\ell+1} = \sigma \left(\sum_{j \in A_i} a_j^\ell w_{ij}^\ell + \beta_i^\ell \right)$$

where A_i is the set of anterior neurons.

The Sigmoid Activation Function

Neural

Networks Pt.
2

W. Guss &
P. Kuznetsov

Feed-Forward
Neural
Networks

How It Works
An
Implementation

Training
Nonconvex
Optimization
Training Demo
Deep Learning




Figure: Two sigmoid functions and the perceptron $\text{step}(x)$.

Definition

We say that $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ is a **sigmoid activation function** if $\sigma(x) \rightarrow 1$ as $x \rightarrow \infty$ and $\sigma(x) \rightarrow -1$ or 0 as $x \rightarrow -\infty$

The Feed-Forward Algorithm

Neural
Networks Pt.
2


W. Guss &
P. Kuznetsov

Feed-Forward
Neural
Networks

How It Works
An
Implementation


Training
Nonconvex
Optimization
Training Demo

Deep Learning



- The feed-forward algorithm propagates input through the neural network layer by layer.
- On every layer each neuron accumulates input from previous layers and then *activates* through the sigmoid activation function.

The Feed-Forward Algorithm



Algorithm 1 An Intuitive Version

```
1: for layer  $\ell = 1$  to  $L - 1$  do
2:   for neuron  $a_j$  on layer  $\ell$  do
3:      $a_j.activate()$ 
4:     for neuron  $a_i$  on layer  $\ell + 1$  do
5:        $a_i.feed(w_{ij}a_j)$ 
6:     end for
7:   end for
8: end for
```

Our Network Implementation

Neural
Networks Pt.
2

W. Guss &
P. Kuznetsov

Feed-Forward
Neural
Networks

How It Works
An
Implementation

Training

Nonconvex
Optimization
Training Demo

Deep Learning

Network

- List of neurons and List of connections
- Feedforward Method
- Backpropagation Method

Neuron

- List of connections anterior and posterior connections
- Feed Method
- Activation Method
- Error update

Our Network Implementation

Neural
Networks Pt.
2

W. Guss &
P. Kuznetsov

Feed-Forward
Neural
Networks

How It Works
An
Implementation

Training

Nonconvex
Optimization
Training Demo

Deep Learning

Connection

General weight information

- Reference to anterior and posterior neuron
- Weight value
- Feedforward

Training

Neural
Networks Pt.
2

W. Guss &
P. Kuznetsov

Feed-Forward
Neural
Networks

How It Works
An
Implementation

Training

Nonconvex
Optimization
Training Demo

Deep Learning

- The goal of machine learning is to adjust learning parameters to better approximate a function.
- For neural networks, these learning parameters are the weights and the bias values.
- We train our network by trying to minimize a loss function:

$$E = \sum (\delta_i - y_i)^2$$

where δ is the expected output and y is the actual output.

- Minimizing the loss function is a nonconvex problem - we need to develop heuristics to properly train the network
- One such heuristic is called gradient descent

Gradient Descent and Error Backpropagation (Calculus Heavy)

Neural

Networks Pt.
2

W. Guss &
P. Kuznetsov

Feed-Forward
Neural
Networks


How It Works
An
Implementation

Training

Nonconvex
Optimization

Training Demo

Deep Learning



- The goal is to travel down the gradient (slope at a point) towards a minima of the error function.
- Error backpropagation is an algorithm that calculates the gradient and then updates the weights according to a *learning rate*.

Definition

We call ∇E the **gradient** of E if

$$\nabla E = \left(\frac{\partial E}{\partial w_{00}^{(0)}}, \frac{\partial E}{\partial w_{01}^{(0)}}, \dots, \frac{\partial E}{\partial w_{ij}^{(L)}} \right).$$

Algorithm 2 The Weight Update Rule

- 1: **for** every weight w_{ij}^ℓ **do**
 - 2: calculate $\Delta w_{ij}^\ell = -\alpha \frac{\partial E}{\partial w_{ij}^\ell}$
 - 3: $w_{ij}^\ell + \Delta w_{ij}^\ell \rightarrow w_{ij}^\ell$
 - 4: **end for**
-

Program Reference

Neural
Networks Pt.
2

W. Guss &
P. Kuznetsov

Feed-Forward
Neural
Networks

How It Works
An
Implementation

Training

Nonconvex
Optimization

Training Demo

Deep Learning

Example (Creating a network, loading a dataset, and training the net)

```
n = Network(<list of layer sizes>)
t = Trainer(n)
t.load_data(<string of the filename>)
t.interactive_step(<learning_rate> [, optional verbose flag])
```

Example (A dataset file)

```
# This is a sample data file
# You can comment with ####
# This dataset has two inputs and one desired output
{0,1} -> {1};
{1,0} -> {1};
{0,0} -> {0.3};
```


Deep Learning

Neural Networks Pt. 2

W. Guss &
P. Kuznetsov

Feed-Forward
Neural
Networks
How It Works
An
Implementation

Training
Nonconvex
Optimization
Training Demo
Deep Learning



- A more powerful variation of the artificial neural network that uses deep architecture, where there are many hidden layers in a network
- Represents problems as a hierarchy of concepts and representations
- Each concept is defined in relation to simpler concepts.

Deep Learning

Neural
Networks Pt.
2

W. Guss &
P. Kuznetsov

Feed-Forward
Neural
Networks

How It Works
An
Implementation

Training
Nonconvex
Optimization
Training Demo

Deep Learning

- Neural networks saw a number of falls from popularity
- Two main events occurred that brought neural networks back
 - The availability of better hardware and more extensive dataset (think big data)
 - The creation of more simplified training methods for deep network architectures
- With the availability of better hardware and large datasets (think big data)

References

Neural
Networks Pt.
2

W. Guss &
P. Kuznetsov

Feed-Forward
Neural
Networks

How It Works
An
Implementation

Training

Nonconvex
Optimization

Training Demo

Deep Learning



Yaniv Taigman et. al. (2014)

DeepFace: Closing the Gap to Human-Level Performance in Face Verification

Facebook AI Research



Michael Nielsen (2014)

Neural Networks and Deep Learning

<http://neuralnetworksanddeeplearning.com/>

Neural
Networks Pt.
2

W. Guss &
P. Kuznetsov

Feed-Forward
Neural
Networks

How It Works
An
Implementation

Training

Nonconvex
Optimization
Training Demo

Deep Learning

The End